



TEXAS ADVANCED COMPUTING CENTER

WWW.TACC.UTEXAS.EDU



TEXAS

The University of Texas at Austin

Tutorial on OpenMP programming

Victor Eijkhout

SSIASC 2016

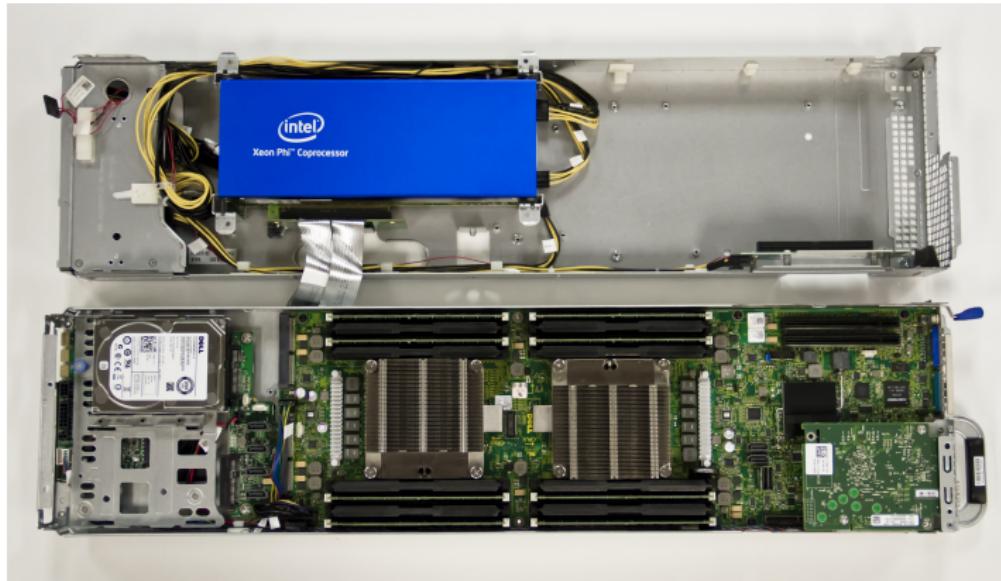
Justification

OpenMP is a flexible tool for incrementally parallelizing a shared memory-based code. This course introduces the main concepts through lecturing and exercises.

The Fork-Join model

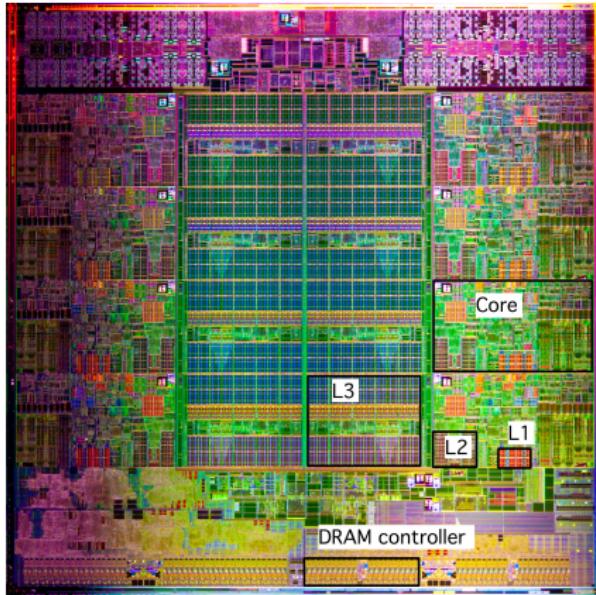
Computer architecture terminology

One cluster node:



A node will have 1 or 2 or (sometimes) 4 'sockets': processor chips.
There may be a co-processor attached.

Structure of a socket

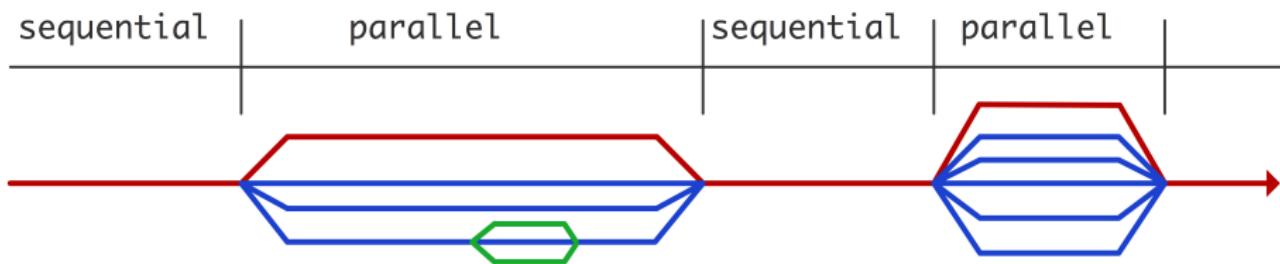


Eight cores per socket, making 16 per node.
They all access the same data.

Threads

Process: stream of instructions

Thread: process can duplicate itself, same code, access to same data



The OS will place threads on different cores: parallel performance.

Note: threads are software. More threads than cores or fewer is allowed.

To write an OpenMP program

```
#include "omp.h"
```

in C, and

```
use omp_lib
```

or

```
#include "omp_lib.h"
```

for Fortran.

To compile an OpenMP program

```
# gcc  
gcc -o foo foo.c -fopenmp  
# Intel compiler  
icc -o foo foo.c -openmp
```

To run an OpenMP program

```
export OMP_NUM_THREADS=8  
./my_omp_program
```

Stampede has 16 cores; more than 16 threads does not make sense.

Quick experiment:

```
for t in 1 2 4 8 12 16; do  
    OMP_NUM_THREADS=$t ./my_omp_program  
done
```

Exercise 1 (parallel)

Write a program that contains the following lines:

```
printf("There are %d processors\n", omp_get_num_procs());  
#pragma omp parallel  
printf("There are %d threads\n",  
    /* !!!! something missing here !!!! */ );
```

The first print statement tells you the number of available cores in the hardware. Your assignment is to supply the missing function that reports the number of threads used. Compile and run the program. Experiment with the OMP_NUM_THREADS environment variable. What do you notice about the number of lines printed?

What happens if I press that button?

Who of you has tried setting the number of threads (much) larger than the number of cores? What happened?

Threads and threads

- Threads are software, cores are hardware.
- The OS can move threads between cores: not a good idea for performance.
- Set: `export OMP_PROC_BIND=true` and you'll be good in most of the cases.
- Look up 'affinity' in the OMP standard for all the details.

Exercise 2

Extend the program from exercise 1. Make a complete program based on these lines:

```
int tsum=0;  
#pragma omp parallel  
    tsum += /* the thread number */  
printf("Sum is %d\n",tsum);
```

Compile and run again. (In fact, run your program a number of times.) Do you see something unexpected? Can you think of an explanation?

Shared memory problems

Race condition: simultaneous update of shared data:

process 1: $I = I + 2$

process 2: $I = I + 3$

Results can be indeterminate:

scenario 1.	scenario 2.	scenario 3.
$I = 0$		
read $I = 0$ compute $I = 2$ write $I = 2$	read $I = 0$ compute $I = 3$ write $I = 3$	read $I = 0$ compute $I = 2$ write $I = 2$
$I = 3$		
$I = 2$		
$I = 5$		

Loop parallelism

Loop parallelism

Much of parallelism in scientific computing is in loops:

- Vector updates and inner products
- Matrix-vector and matrix-matrix operations
- Finite Element meshes
- Multigrid

Work distribution

- Suppose loop iterations are independent:
- Distribute them over the threads:
- Use `omp_get_thread_num` to determine disjoint subsets.
- How would you do this specifically?

Workshare constructs

Here's the two-step parallelization in OpenMP:

- You use the `parallel` directive to create a team of threads;
- then you use a 'workshare' construct to distribute the work over the team.
- For loops that is the `for` (or `do`) construct.

Workshare construct for loops

C: directive followed by statement or block:

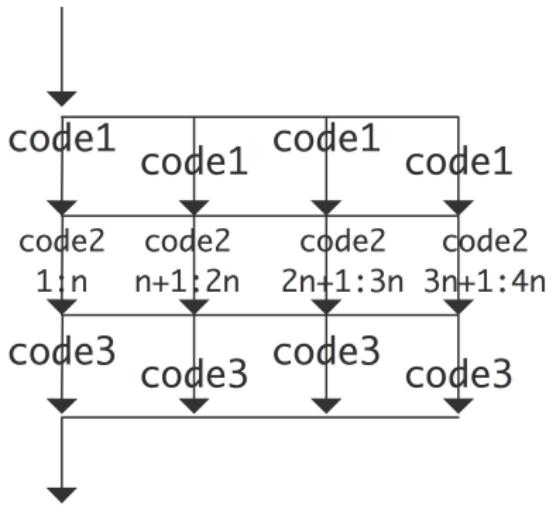
```
#pragma omp parallel
{
#pragma omp for
    for (i=0; i<N; i++)
        ... something with i ...
}
```

Fortran: matching end directive

```
!$omp parallel
!$omp do
    do i=1,n
        ... something with i ...
    end do
!$omp end do
!$omp end parallel
```

Stuff inside a parallel region

```
#pragma omp parallel
{
    code1();
#pragma omp for
    for (i=1; i<=4*N; i++) {
        code2();
    }
    code3();
}
```



Exercise 3 (pi)

Compute π by *numerical integration*. We use the fact that π is the area of the unit circle, and we approximate this by computing the area of a quarter circle using *Riemann sums*.

- Let $f(x) = \sqrt{1 - x^2}$ be the function that describes the quarter circle for $x = 0 \dots 1$;
- Then we compute

$$\pi/4 \approx \sum_{i=0}^{N-1} \Delta x f(x_i) \quad \text{where } x_i = i\Delta x \text{ and } \Delta x = 1/N$$

Write a program for this, and parallelize it using OpenMP parallel for directives.

- Put a parallel directive around your loop. Does it still compute the right result? Does the time go down with the number of threads? (The answers should be no and no.)
- Change the parallel to parallel for (or parallel do). Now is the result correct? Does execution speed up? (The answers should now be no and yes.)
- Put a critical directive in front of the update. (Yes and very much no.)
- Remove the critical and add a clause reduction (+:quarterpi) to the for directive. Now it should be correct and efficient.

Use different numbers of cores and compute the speedup you attain over the sequential computation. Is there a performance difference between the C
the sequential code?

Loop schedules

- Default: static scheduling of iterations.

Very efficient. Good if all iterations take the same amount of time.
schedule (static)

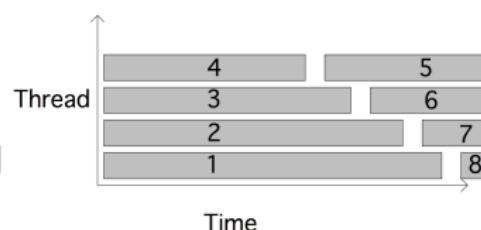
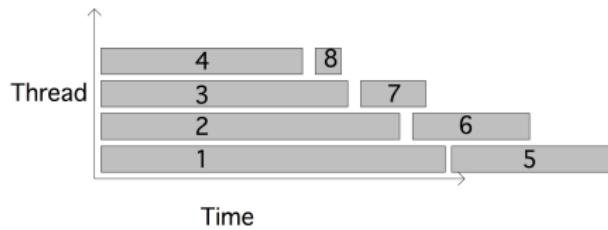
- Other possibility: dynamic.

Runtime overhead; better if iterations do not take the same amount of time.

schedule (dynamic)

Four threads, 8 tasks of decreasing size

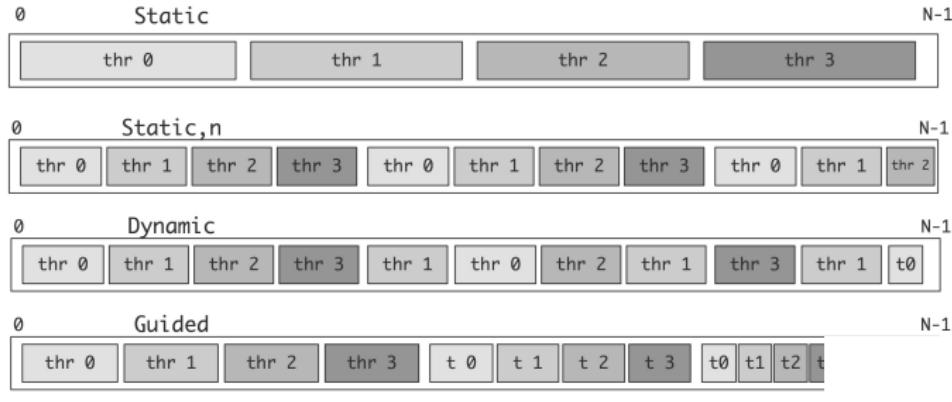
dynamic schedule is better:



Chunk size

With N iterations and t threads:

- Static: each thread gets N/t iterations.
explicit chunk size: `schedule(static, 123)`
- Dynamic: each thread gets 1 iteration at a time
explicit chunk size: `schedule(dynamic, 45)`
- Help from OpenMP:
guided schedule uses decreasing chunk size (with optional minimum chunk):
`schedule(guided, 6)`



Reductions

- Inner product loop:

```
s = 0.;  
for (i=0; i<N; i++)  
    s += x[i]*y[i];
```

- Use the reduction(+:s) clause.
- All the usual operations are available; you can also make your own.

Exercise 4 (piadapt)

We continue with exercise 3. We add ‘adaptive integration’: where needed, the program refines the step size¹. This means that the iterations no longer take a predictable amount of time.

```
for (i=0; i<nsteps; i++) {  
    double  
        x = i*h, x2 = (i+1)*h,  
        y = sqrt(1-x*x), y2 = sqrt(1-x2*x2),  
        slope = (y-y2)/h;  
    if (slope>15) slope = 15;  
    int  
        samples = 1+(int)slope, is;  
    for (is=0; is<samples; is++) {  
        double  
            hs = h/samples,  
            xs = x+ is*hs,  
            ys = sqrt(1-xs*xs);  
        quarterpi += hs*ys;  
        nsamples++;  
    }  
}  
pi = 4*quarterpi;
```

same exercise

- ① Use the `omp parallel for` construct to parallelize the loop. As in the previous lab, you may at first see an incorrect result. Use the `reduction` clause to fix this.
- ② Your code should now see a decent speedup, using up to 8 cores. However, it is possible to get completely linear speedup. For this you need to adjust the schedule.
Start by using `schedule(static, n)`. Experiment with values for n . When can you get a better speedup? Explain this.
- ③ Since this code is somewhat dynamic, try `schedule(dynamic)`. This will actually give a fairly bad result. Why? Use `schedule(dynamic, n)` instead, and experiment with values for n .
- ④ Finally, use `schedule(guided)`, where OpenMP uses a heuristic. What results does that give?
- ⑤ `schedule(auto)` : leave it up to the system.
- ⑥ `schedule(runtime)` : leave it up to environment variables; good for experimenting.

More loop topics

- Multiple loops can be collapsed: `collapse(2)`. Improves performance.
- Ordered iterations: normally OpenMP can execute iterations in any sequence. You can force ordering if you absolutely have to. Bad for performance!
- There is a barrier at the end of a `for`: use `nowait` to let threads continue.

Workshare constructs

What is worksharing again?

- The `omp parallel` creates a team of threads.
- Now you need to distributed work among them.
- Already seen: `for`, `do`
- Similar: `sections`
- Not obvious: `single`
- Fortran only: `workshare` (works with array notation, but compiler support seems mediocre)
- Story in itself: `task`

Sections

Independent separate calculations:

```
double fx = f(x), gx = g(x), hx = h(x);  
..... fx ... gx ... hx .....
```

```
#pragma omp sections  
{  
    #pragma omp section  
    fx = f(x)  
    #pragma omp section  
    gx = g(x)  
    #pragma omp section  
    hx = h(x)  
}
```

Adding them together:

```
s = f(x)+g(x)+h(x);
```

Use reduction.

Single

```
int a;  
#pragma omp parallel  
{  
    #pragma omp single  
    a = f(); // some computation  
    #pragma omp sections  
    // various different computations using 'a'  
}
```

- Is executed by a single thread.
- Has implicit barrier, so the result is available to everyone after.
- `master` is similar, does not have barrier.

Exercise 5

What is the difference between this approach and how the same computation would be parallelized in MPI?

Thread data

Shared and private data

You have already seen some of the basics:

- Data declared outside a parallel region is shared.
- Data declared in the parallel region is private.
(Fortran does not have this block scope mechanism)

```
int i;  
#pragma omp parallel  
{ double i; .... }
```

- You can change all this with clauses:

```
int i;  
#pragma omp parallel private(i)
```

Variables in loops

```
int i; double t;  
#pragma omp parallel for  
for (i=0; i<N; i++) {  
    t = sin(i*pi*h);  
    x[i] = t*t;  
}
```

- The loop variable is automatically private.
- The temporary `t` is shared, but conceptually private to each iteration:
needs to be declared private.
(What happens if you don't?)

Copying to/from private data

- Private data is uninitialized

```
int i = 3;  
#pragma omp parallel private(i)  
    printf("%d\n", i); // undefined!
```

- To import a value:

```
int i = 3;  
#pragma omp parallel firstprivate(i)  
    printf("%d\n", i); // undefined!
```

- lastprivate to preserve value of last iteration.

Default behaviour

- `default(shared)` or `default(private)`
- useful for debugging: `default(none)`
because you have to specify everything as shared/private

Persistent thread data

- Private data disappears after the parallel region.

What if you want data to persist?

- Directive `threadprivate`

```
double seed;  
#pragma omp threadprivate(seed)
```

- Standard application: random number generation.
- Tricky: has to be global or static.

Arrays

- Statically allocated arrays can be made private.
- Dynamically allocated ones can not: the pointer becomes private.

Synchronization

Need for synchronization

- The loop and sections directives do not specify an ordering, sometimes you want to force an ordering.
- Barriers: global synchronization.
- Critical sections: only one process can execute a statement this prevents race conditions.
- Locks: protect data items from being accessed.

Barriers

- Every workshare construct has an implicit barrier:

```
#pragma omp parallel
{
    #pragma omp for
    for ( .. i .. )
        x[i] = ...
    #pragma omp for
    for ( .. i .. )
        y[i] = .. x[i] .. x[i+1] .. x[i-1] ...
}
```

First loop is completely finished before second.

- Explicit barrier:

```
#pragma omp parallel
{
    x = f();
#pragma omp barrier
    .... x ....
```

Critical sections

- Critical section: One update at a time.

```
#pragma omp parallel
{
    double x = f();
#pragma omp critical
    global_update(x);
}
```

- `atomic`: special case for simple operations, possible hardware support

```
#pragma omp atomic
    t += x;
```

Warning

- Critical sections are not cheap! The operating system takes thousands of cycles to coordinate the threads.
- Use only if minor amount of work.
- Do not use if a reduction suffices.
- Name your critical sections.
- Explore locks if there may not be a data conflict.

Locks

- Critical sections are coarse:
they dictate exclusive access to a *statement*
- Suppose you update a big table
updates to non-conflicting locations should be allowed
- Locks protect a single data item.

Tasks

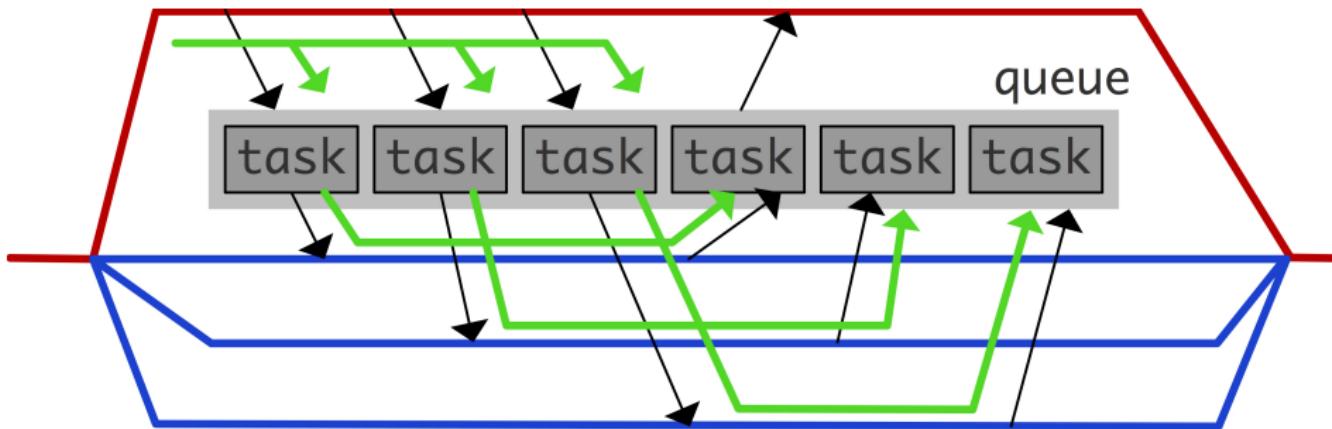
More flexibility

- You have seen loops and sections.
- How about linked lists or trees?
- Tasks are very flexible:
you create work, it goes on a queue, gets executed later

```
p = head_of_list();
while (!end_of_list(p)) {
#pragma omp task
    process( p );
    p = next_element(p);
}
```

Threads, tasks, queues

- There is one queue (per team), not visible to the programmer.
- One thread starts generating tasks.
- Tasks can recursively generate tasks.
- You never know who executes what.



Exercise 6 (taskfactor)

Use tasks to find the smallest factor of a large number (using $2999 \cdot 3001$ as test case): generate a task for each trial factor. Start with this code:

```
int factor=0;  
#pragma omp parallel  
#pragma omp single  
for (int f=2; f<4000; f++) {  
    { // see if 'f' is a factor  
        if (N%f==0) { // found factor!  
            factor = f;  
        }  
    }  
    if (factor>0)  
        break;  
}  
if (factor>0)  
    printf("Found a factor: %d\n", facto
```

Task synchronization

Mechanisms for task synchronization:

- `taskwait`: wait for all previous tasks (not nested)
- `taskgroup`: wait for all tasks, including nested
- `depend`: synchronize on data items.

Example: tree traversal

```
int process( node n ) {  
    if (n.is_leaf)  
        return n.value;  
    for ( c : n.children) {  
#pragma omp task  
        process(c);  
#pragma omp taskwait  
        return sum_of_children();  
    }  
}
```

Example: Fibonacci

```
long fib(int n) {  
    if (n<2) return n;  
    else { long f1,f2;  
#pragma omp task  
    f1 = fib(n-1);  
#pragma omp task  
    f2 = fib(n-2);  
#pragma omp taskwait  
    return f1+f2;  
}  
  
#pragma omp parallel  
#pragma omp single  
printf("Fib(50)=%ld",fib(50));
```

(what is conceptually wrong with this example?)

Fibonacci once more

```
long fibs[100];
void fib(n) {
    if (n>=2) {
        #pragma omp task \
            depend( in:fibs[n-2],in:fibs[n-1] ) \
            depend( out:fibs[n] )
        fibs[n] = fibs[n-2]+fibs[n-1];
    }
}

#pragma omp parallel
#pragma omp single
for (i<50)
    fib(i);
```

Remaining topics

Affinity

How do you place threads on cores?

- Two socket design NUMA
- Intel KNL has quadrants and hardware multi-threading
- `OMP_PROC_BIND` and `OMP_PLACES`

Accelerators

OpenMP 4 has mechanisms for offloading.

SIMD

Processors have 4 or 8-wide SIMD.
convert OpenMP loop to SIMD vector instructions.

Index

- atomic, 43
- collapse(2), 27
- default(none), 37
- default(private), 37
- default(shared), 37
- depend, 50
- do, 18, 29
- for, 18, 27, 29
- lastprivate, 36
- master, 31
- nowait, 27
- numerical integration, 21
- omp parallel, 29
- omp parallel for, 26
- omp_get_thread_num, 17
- OMP_PLACES, 55
- OMP_PROC_BIND, 55
- parallel, 18
- reduction, 26
- reduction(+:s), 24
- **Riemann sums**, 21
- schedule(auto), 26
- schedule(dynamic), 26
- schedule(dynamic, n), 26
- schedule(guided), 26
- schedule(runtime), 26
- schedule(static, n), 26
- sections, 29
- single, 29
- task, 29
- taskgroup, 50
- taskwait, 50
- threadprivate, 38
- workshare, 29