

Parallelism in Applied Math Calculations

Victor Eijkhout

374/394 spring 2013

ODEs and PDEs

Time-evolving phenomena: IVP (Initial Value Problem), usually Ordinary Differential Equations

Space-constraint phenomena: BVP (Boundary Value Problem), usually Partial Differential Equations

Approximating Differential Equations

Finite difference approximation

We turn the continuous problem into a discrete one, by looking at finite time/space steps.

Assume all functions are sufficiently smooth, and use Taylor series:

$$u(t + \Delta t) = u(t) + u'(t)\Delta t + u''(t)\frac{\Delta t^2}{2!} + u'''(t)\frac{\Delta t^3}{3!} + \dots$$

This gives for u' :

$$u'(t) = \frac{u(t + \Delta t) - u(t)}{\Delta t} + O(\Delta t^2)$$

So we approximate

$$u'(t) \approx \frac{u(t + \Delta t) - u(t)}{\Delta t}$$

and the “truncation error” is $O(\Delta t^2)$.

Finite differences 2

How does this help? In $u' = f(t, u)$ substitute

$$u'(t) \rightarrow \frac{u(t + \Delta t) - u(t)}{\Delta t}$$

giving

$$\frac{u(t + \Delta t) - u(t)}{\Delta t} = f(t, u(t))$$

or

$$u(t + \Delta t) = u(t) + \Delta t f(t, u(t))$$

Let $t_0 = 0, t_{k+1} = t_k + \Delta t = \dots = (k + 1)\Delta t, u(t_k) = u_k$:

$$u_{k+1} = u_k + \Delta t f(t_k, u_k)$$

Discretization

'Explicit Euler' or 'Euler forward'.

Does this compute something close to the true solution?

'Discretization error'

Boundary value problems

Consider $u''(x) = f(x, u, u')$ for $x \in [a, b]$ where $u(a) = u_a$, $u(b) = u_b$ in 1D and

$$-u_{xx}(\bar{x}) - u_{yy}(\bar{x}) = f(\bar{x}) \text{ for } x \in \Omega = [0, 1]^2 \text{ with } u(\bar{x}) = u_0 \text{ on } \delta\Omega. \quad (1)$$

in 2D.

Approximation of 2nd order derivatives

Taylor series (write h for δx):

$$u(x+h) = u(x) + u'(x)h + u''(x)\frac{h^2}{2!} + u'''(x)\frac{h^3}{3!} + u^{(4)}(x)\frac{h^4}{4!} + u^{(5)}(x)\frac{h^5}{5!} + \dots$$

and

$$u(x-h) = u(x) - u'(x)h + u''(x)\frac{h^2}{2!} - u'''(x)\frac{h^3}{3!} + u^{(4)}(x)\frac{h^4}{4!} - u^{(5)}(x)\frac{h^5}{5!} + \dots$$

Subtract:

$$u(x+h) + u(x-h) = 2u(x) + u''(x)h^2 + u^{(4)}(x)\frac{h^4}{12} + \dots$$

so

$$u''(x) = \frac{u(x+h) - 2u(x) + u(x-h)}{h^2} - u^{(4)}(x)\frac{h^4}{12} + \dots$$

Numerical scheme:

$$-\frac{u(x+h) - 2u(x) + u(x-h)}{h^2} = f(x, u(x), u'(x))$$

(2nd order PDEs are very common!)

This leads to linear algebra

$$-\frac{u(x+h) - 2u(x) + u(x-h)}{h^2} = f(x, u(x), u'(x))$$

Equally spaced points on $[0, 1]$: $x_k = kh$ where $h = 1/n$, then

$$-u_{k+1} + 2u_k - u_{k-1} = h^2 f(x_k, u_k, u'_k) \quad \text{for } k = 1, \dots, n-1$$

Written as matrix equation:

$$\begin{pmatrix} 2 & -1 & & \\ -1 & 2 & -1 & \\ & \ddots & \ddots & \ddots \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \\ \vdots \end{pmatrix} = \begin{pmatrix} h^2 f_1 + u_0 \\ h^2 f_2 \\ \vdots \end{pmatrix}$$

Matrix properties

- Very sparse, banded
- Symmetric (only because 2nd order problem)
- Sign pattern: positive diagonal, nonpositive off-diagonal
(true for many second order methods)
- Positive definite (just like the continuous problem)

Initial Boundary value problem

Heat conduction in a rod $T(x, t)$ for $x \in [a, b]$, $t > 0$:

$$\frac{\partial}{\partial t} T(x, t) - \alpha \frac{\partial^2}{\partial x^2} T(x, t) = q(x, t)$$

- Initial condition: $T(x, 0) = T_0(x)$
- Boundary conditions: $T(a, t) = T_a(t)$, $T(b, t) = T_b(t)$
- Material property: $\alpha > 0$ is thermal diffusivity
- Forcing function: $q(x, t)$ is externally applied heating.

The equation $u''(x) = f$ above is the steady state.

Discretization

Space discretization: $x_0 = a$, $x_n = b$, $x_{j+1} = x_j + \Delta x$

Time discretiation: $t_0 = 0$, $t_{k+1} = t_k + \Delta t$

Let T_j^k approximate $T(x_j, t_k)$

Space:

$$\frac{\partial}{\partial t} T(x_j, t) - \alpha \frac{T(x_{j-1}, t) - 2T(x_j, t) + T(x_{j+1}, t)}{\Delta x^2} = q(x_j, t)$$

Explicit time stepping:

$$\frac{T_j^{k+1} - T_j^k}{\Delta t} - \alpha \frac{T_{j-1}^k - 2T_j^k + T_{j+1}^k}{\Delta x^2} = q_j^k$$

Implicit time stepping:

$$\frac{T_j^{k+1} - T_j^k}{\Delta t} - \alpha \frac{T_{j-1}^{k+1} - 2T_j^{k+1} + T_{j+1}^{k+1}}{\Delta x^2} = q_j^{k+1}$$

Computational form: explicit

$$T_j^{k+1} = T_j^k + \frac{\alpha \Delta t}{\Delta x^2} (T_{j-1}^k - 2T_j^k + T_{j+1}^k) + \Delta t q_j^k$$

This has an explicit form:

$$\mathcal{T}^{k+1} = \left(I + \frac{\alpha \Delta t}{\Delta x^2} \right) \mathcal{T}^k + \Delta t \mathcal{q}^k$$

Computational form: implicit

$$T_j^{k+1} - \frac{\alpha \Delta t}{\Delta x^2} (T_{j-1}^k - 2T_j^k + T_{j+1}^k) = T_j^k + \Delta t q_j^k$$

This has an implicit form:

$$\left(I - \frac{\alpha \Delta t}{\Delta x^2} K \right) \tilde{T}^{k+1} = \tilde{T}^k + \Delta t \tilde{q}^k$$

Needs to solve a linear system in every time step

Stability of explicit scheme

Needs

$$\Delta t < \frac{\Delta x^2}{2\alpha}$$

big restriction on size of time steps

Stability of implicit scheme

Stability condition always satisfied: method always stable.

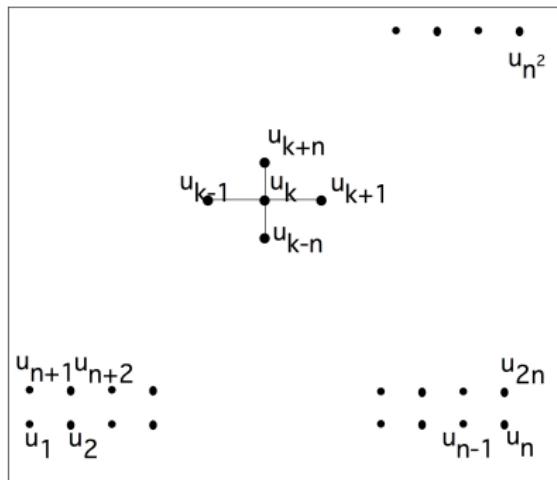
Sparse matrix in 2D case

Sparse matrices so far were tridiagonal: only in 1D case.

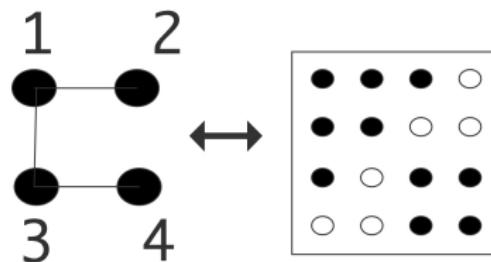
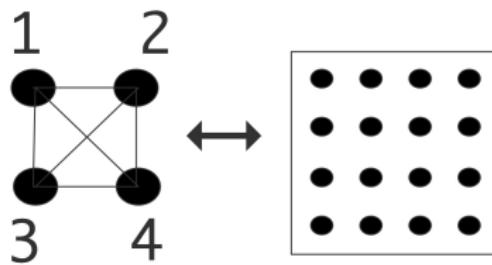
Two-dimensional: $-u_{xx} - u_{yy} = f$ on unit square $[0, 1]^2$

Difference equation:

$$4u(x, y) - u(x + h, y) - u(x - h, y) - u(x, y + h) - u(x, y - h) = h^2 f(x, y)$$



Graph theory of sparse matrices



The graph view of things

Poisson eq:

$$4u_k - u_{k-1} - u_{k+1} - u_{k-n} - u_{k+n} = f_k$$

Consider a graph where $\{u_k\}_k$ are the edges
and (u_i, u_j) is an edge iff $a_{ij} \neq 0$.

This is the (adjacency) graph of a sparse matrix.

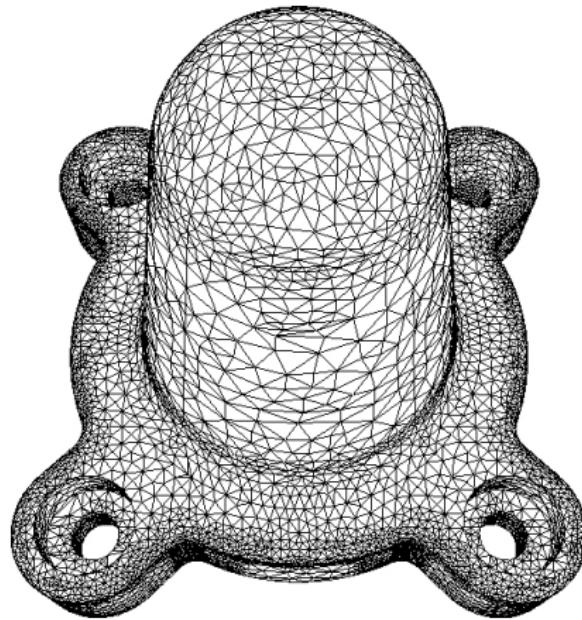
Sparse matrix from 2D equation

$$\left(\begin{array}{cccc|ccc|c} 4 & -1 & & \emptyset & -1 & & \emptyset & \\ -1 & 4 & 1 & & & -1 & & \\ \ddots & \ddots & \ddots & & & \ddots & & \\ & \ddots & \ddots & -1 & & \ddots & & \\ \emptyset & & -1 & 4 & \emptyset & & -1 & \\ \hline -1 & & \emptyset & & 4 & -1 & & -1 \\ & -1 & & & -1 & 4 & -1 & -1 \\ & \uparrow & \ddots & & \uparrow & \uparrow & \uparrow & \uparrow \\ k-n & & & & k-1 & k & k+1 & k+n \\ \hline & & & -1 & & & -1 & 4 \\ & & & & \ddots & & \ddots & \ddots \end{array} \right)$$

Matrix properties

- Very sparse, banded
- Symmetric (only because 2nd order problem)
- Sign pattern: positive diagonal, nonpositive off-diagonal
(true for many second order methods)
- Positive definite (just like the continuous problem)
- Constant diagonals: only because of the constant coefficient differential equation

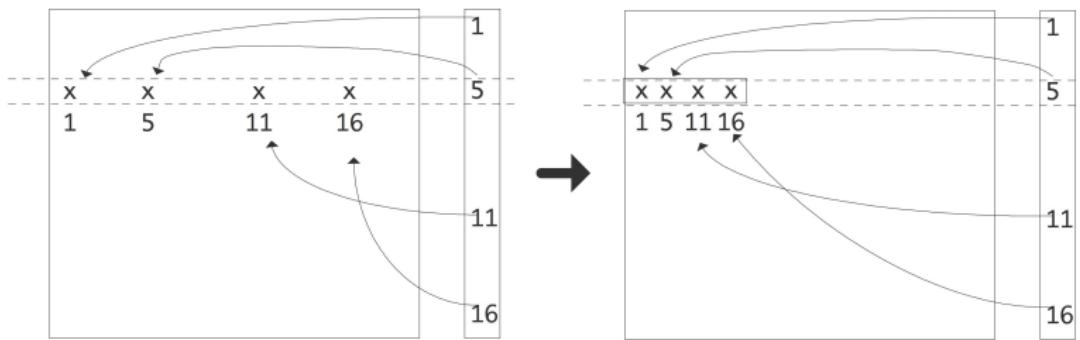
Realistic meshes



Sparse matrix storage

Matrix above has many zeros: n^2 elements but only $O(n)$ nonzeros. Big waste of space to store this as square array.

Matrix is called 'sparse' if there are enough zeros to make specialized storage feasible.



Compressed Row Storage

$$A = \begin{pmatrix} 10 & 0 & 0 & 0 & -2 & 0 \\ 3 & 9 & 0 & 0 & 0 & 3 \\ 0 & 7 & 8 & 7 & 0 & 0 \\ 3 & 0 & 8 & 7 & 5 & 0 \\ 0 & 8 & 0 & 9 & 9 & 13 \\ 0 & 4 & 0 & 0 & 2 & -1 \end{pmatrix}. \quad (2)$$

Compressed Row Storage (CRS): store all nonzeros by row, their column indices, pointers to where the columns start (1-based indexing):

val	10	-2	3	9	3	7	8	7	3	...	9	13	4	2	-1
col_ind	1	5	1	2	6	2	3	4	1	...	5	6	2	5	6
row_ptr	1	3	6	9	13	17	20								

Sparse matrix operations

Most common operation: matrix-vector product

```
for (row=0; row<nrows; row++) {  
    s = 0;  
    for (icol=ptr[row]; icol<ptr[row+1]; icol++) {  
        int col = ind[icol];  
        s += a[aptr] * x[col];  
        aptr++;  
    }  
    y[row] = s;  
}
```

Operations with changes to the nonzero structure are much harder!

Indirect addressing of x gives low spatial and temporal locality.

Dense matrix-vector product

Parallel matrix-vector product; dense

- Assume a division by block rows
- Every processor p has a set of row indices I_p

Mvp on processor p :

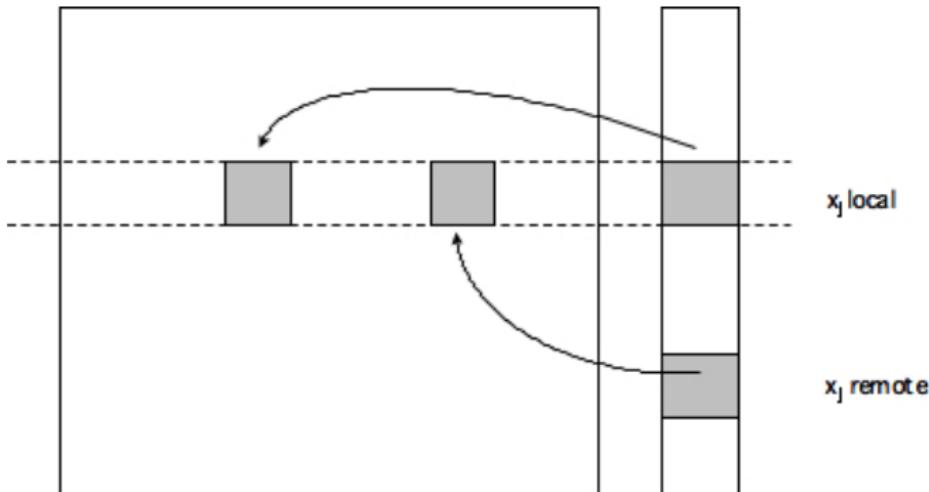
$$\forall_i: y_i = \sum_j a_{ij}x_j$$

$$\forall_i: y_i = \sum_q \sum_{j \in I_q} a_{ij}x_j$$

Local and remote parts:

$$\forall_i: y_i = \sum_{j \in I_p} a_{ij}x_j + \sum_{q \neq p} \sum_{j \in I_q} a_{qj}x_j$$

Local part I_p can be executed right away, I_q requires communication.

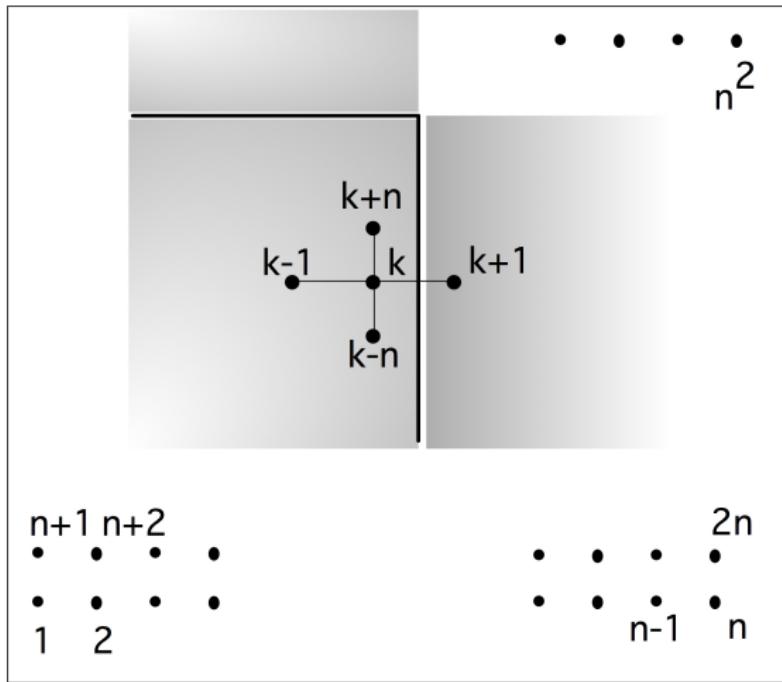


Combine:

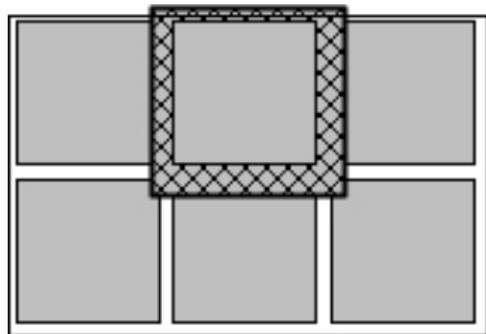
Sparse matrix-vector product

Sparse matrix-vector product

Difference stencil



induces ghost region:



Limited number of neighbours, limited buffer space

Scaling

Separately 1D and 2D partitioning of the domain.

Nested dissection

Fill-in during LU

Fill-in: index (i, j) where $a_{ij} = 0$ but $\ell_{ij} \neq 0$ or $u_{ij} \neq 0$.

2D BVP: Ω is $n \times n$, gives matrix of size $N = n^2$, with bandwidth n .

Matrix storage $O(N)$

LU storage $O(N^{3/2})$

LU factorization work $O(N^2)$

Cute fact: storage can be computed linear in #nonzeros

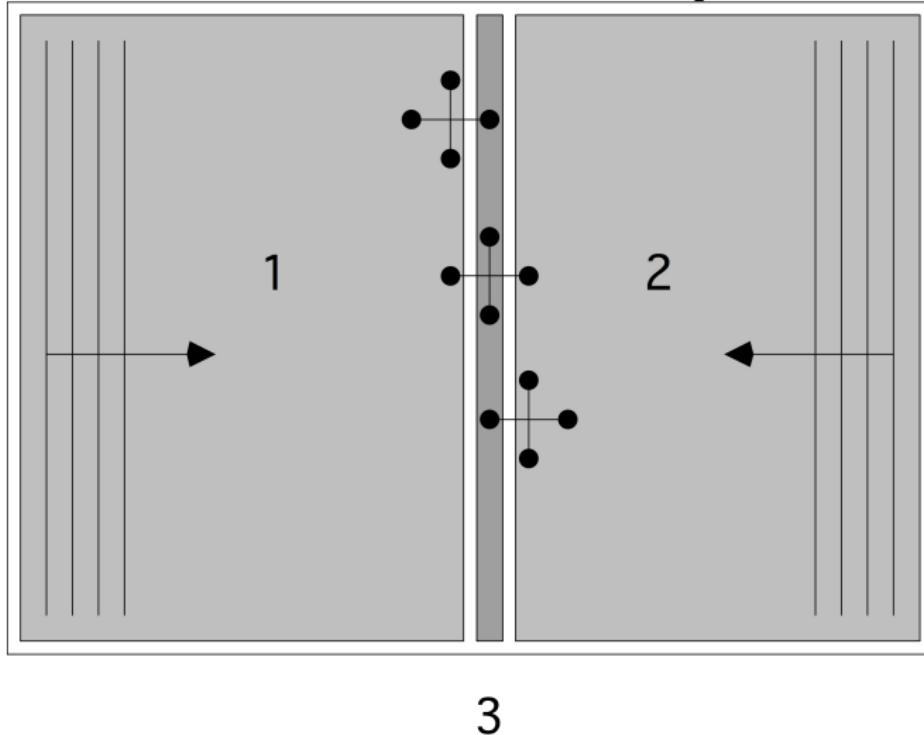
Fill-in is a function of ordering

$$\begin{pmatrix} * & * & \dots & * \\ * & * & & \emptyset \\ \vdots & & \ddots & \\ * & \emptyset & & * \end{pmatrix}$$

After factorization the matrix is dense.

Can this be permuted?

Domain decomposition



$$\left(\begin{array}{c|ccccc|ccccc} \star & \star & & & & & & 0 & \\ \star & \star & \star & & & & & \vdots & \\ \cdot & \cdot & \cdot & \cdot & & & \emptyset & & \\ & \star & \star & \star & & & & 0 & \\ & \star & \star & & & & & \star & \\ \hline & & & \star & \star & & & 0 & \\ & & & \star & \star & \star & & \vdots & \\ & & & \cdot & \cdot & \cdot & & \vdots & \\ & & & \star & \star & \star & & 0 & \\ & & & \star & \star & \star & & \star & \\ \hline 0 & \dots & \dots & 0 & \star & 0 & \dots & \dots & 0 & \star & \star \end{array} \right) \quad \left. \begin{array}{l} (n^2 - n)/2 \\ (n^2 - n)/2 \\ n \end{array} \right\}$$

DD factorization

$$A^{\text{DD}} = \begin{pmatrix} A_{11} & \emptyset & A_{13} \\ \emptyset & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix} =$$
$$\begin{pmatrix} I & & \\ \emptyset & I & \\ A_{31}A_{11}^{-1} & A_{32}A_{22}^{-1} & I \end{pmatrix} \begin{pmatrix} A_{11} & \emptyset & A_{13} \\ A_{22} & A_{23} & \\ & & S \end{pmatrix}$$

$$S = A_{33} - A_{31}A_{11}^{-1}A_{13} - A_{32}A_{22}^{-1}A_{23}$$

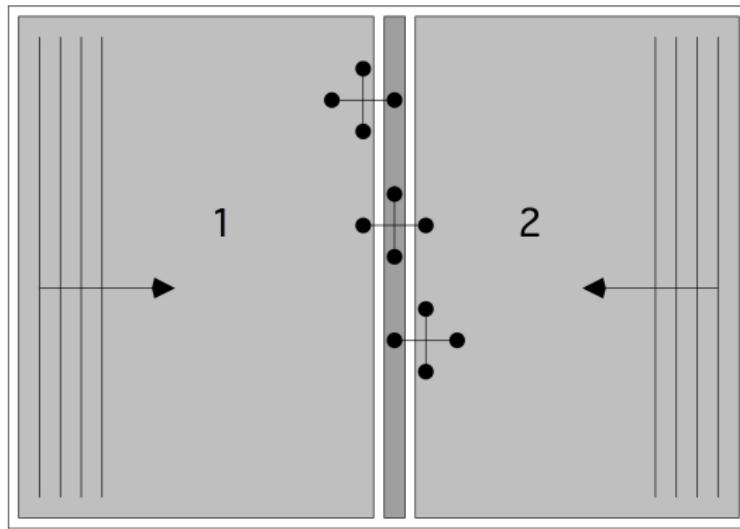
Parallelism . . .

Graph theory of sparse elimination

$$a_{ij} \leftarrow a_{ij} - a_{ik} a_{kk}^{-1} a_{kj}$$



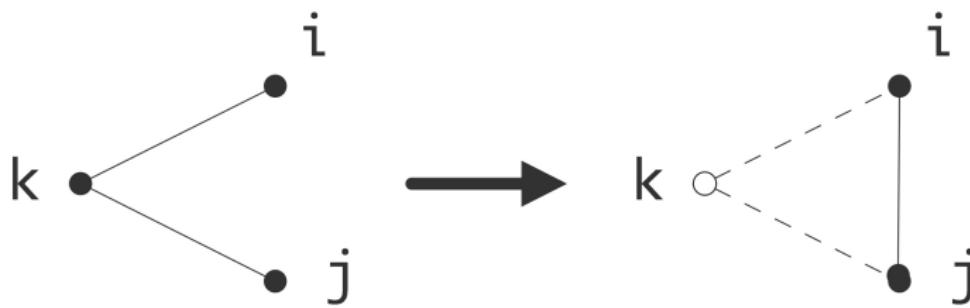
Graph theory of sparse elimination



3

Graph theory of sparse elimination

$$a_{ij} \leftarrow a_{ij} - a_{ik} a_{kk}^{-1} a_{kj}$$



So inductively S is dense

Recursive bisection

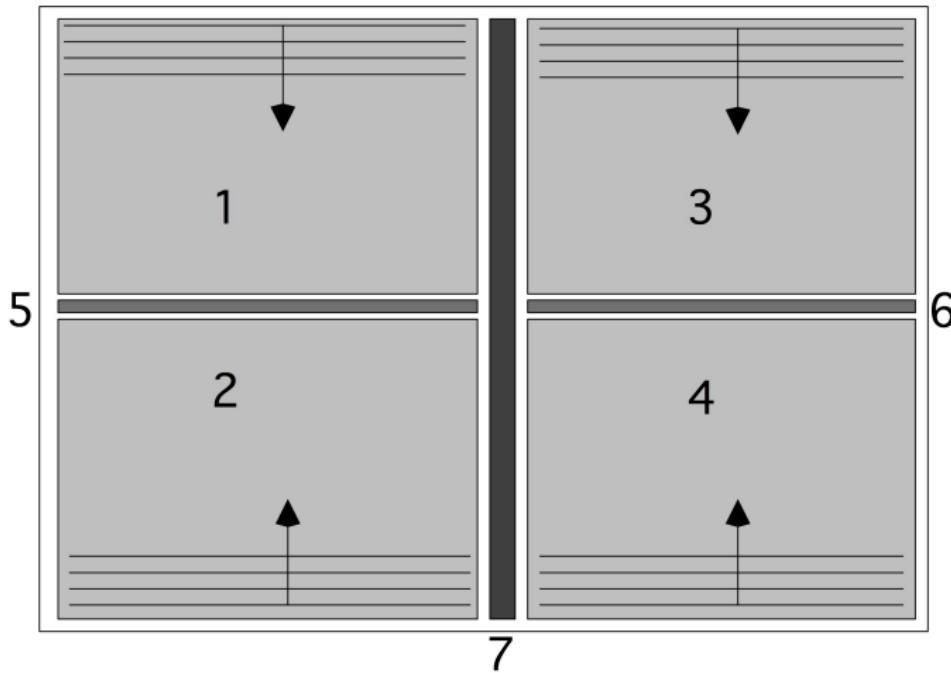
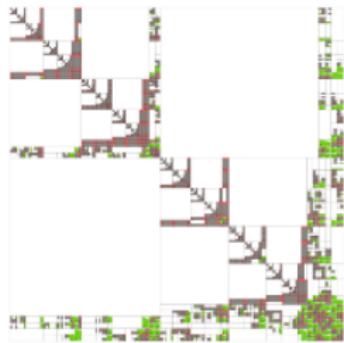


Figure : A four-way domain decomposition

$$A^{\text{DD}} = \begin{pmatrix} A_{11} & & & A_{15} & A_{17} \\ & A_{22} & & A_{25} & A_{27} \\ & & A_{33} & & A_{36} & A_{37} \\ & & & A_{44} & & A_{46} & A_{47} \\ A_{51} & A_{52} & & & A_{55} & & A_{57} \\ & & A_{63} & A_{64} & & A_{66} & A_{67} \\ A_{71} & A_{72} & A_{73} & A_{74} & A_{75} & A_{76} & A_{77} \end{pmatrix}$$

The domain/operator/graph view is more insightful, don't you think?

How does this look in reality?



Complexity

With $n = \sqrt{N}$:

- one dense matrix on a separator of size n , plus
- two dense matrices on separators of size $n/2$
- $\rightarrow 3/2 n^2$ space and $5/12 n^3$ time
- and then four times the above with $n \rightarrow n/2$

$$\begin{aligned}\text{space} &= 3/2n^2 + 4 \cdot 3/2(n/2)^2 + \dots \\ &= N(3/2 + 3/2 + \dots) \quad \log n \text{ terms} \\ &= O(N \log N)\end{aligned}$$

$$\begin{aligned}\text{time} &= 5/12n^3/3 + 4 \cdot 5/12(n/2)^3/3 + \dots \\ &= 5/12N^{3/2}(1 + 1/4 + 1/16 + \dots) \\ &= O(N^{3/2})\end{aligned}$$

More direct factorizations

Minimum degree, multifrontal, . . .

Finding good separators and domain decompositions is tough in general.

Parallel preconditioners

Sparse operations in parallel: mvp

Mvp $y = Ax$

```
for i=1..n  
    y[i] = sum over j=1..n a[i,j]*x[j]
```

In parallel:

```
for i=myfirstrow..mylastrow  
    y[i] = sum over j=1..n a[i,j]*x[j]
```

How about ILU solve?

Consider $Lx = y$

```
for i=1..n
    x[i] = (y[i] - sum over j=1..i-1 ell[i,j]*x[j])
            / a[i,i]
```

Parallel code:

```
for i=myfirstrow..mylastrow
    x[i] = (y[i] - sum over j=1..i-1 ell[i,j]*x[j])
            / a[i,i]
```

Problems?

Block method

```
for i=myfirstrow..mylastrow  
    x[i] = (y[i] - sum over j=myfirstrow..i-1 ell[i,j]*x[j]  
            / a[i,i]
```

Block Jacobi with local GS solve

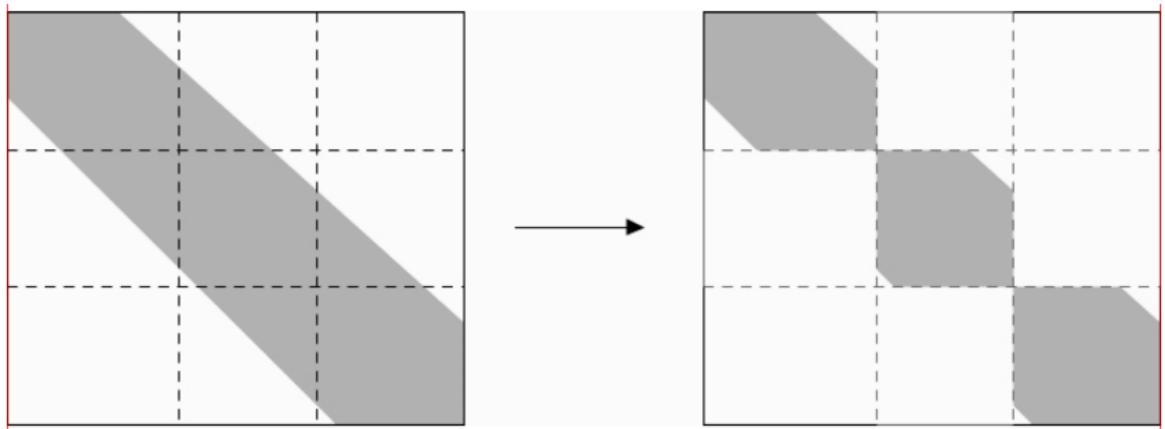


Figure : Sparsity pattern corresponding to a block Jacobi preconditioner

Multicolour ILU

